

Class10

Bayah Essayem (A17303992)

2025-02-06

Table of contents

The PDB database	1
3. Intro to Bio3D in R	10
4. Predicting functional dynamics	12

The PDB database

The main repository of biomolecular structure data is called the PDB found at <https://www.rcsb.org>

Let's see what this database contains. I went to PDB > Analysis > PDB Statistics > By Exp method and molecular type

```
pdstats <- read.csv("Data Export Summary.csv")
pdstats
```

	Molecular.Type	X.ray	EM	NMR	Multiple.methods	Neutron	Other
1	Protein (only)	169,563	16,774	12,578	208	81	32
2	Protein/Oligosaccharide	9,939	2,839	34	8	2	0
3	Protein/NA	8,801	5,062	286	7	0	0
4	Nucleic acid (only)	2,890	151	1,521	14	3	1
5	Other	170	10	33	0	0	0
6	Oligosaccharide (only)	11	0	6	1	0	4
	Total						
1	199,236						
2	12,822						
3	14,156						
4	4,580						

```
5      213
6       22
```

Q1: What percentage of structures in the PDB are solved by X-Ray and Electron Microscopy. A1: percentage solved by x-ray is 82.83549%, while the Electron Microscopy is 10.75017%

```
pdbstats$X.ray
```

```
[1] "169,563" "9,939"  "8,801"  "2,890"  "170"    "11"
```

The comma in these numbers is causing them to be read as characters rather than numeric. This can be fixed by replacing “,” with ‘sub()’ function:

```
x <- pdbstats$X.ray
sum( as.numeric(sub(",", "", x)))
```

```
[1] 191374
```

Or I can use the **readr** package and the ‘read.csv’ function

```
library(readr)
pdbstats <- read_csv("Data Export Summary.csv")
```

```
Rows: 6 Columns: 8
```

```
-- Column specification -----
```

```
Delimiter: ","
```

```
chr (1): Molecular Type
```

```
dbl (3): Multiple methods, Neutron, Other
```

```
num (4): X-ray, EM, NMR, Total
```

```
i Use `spec()` to retrieve the full column specification for this data.
```

```
i Specify the column types or set `show_col_types = FALSE` to quiet this message.
```

```
pdbstats
```

```
# A tibble: 6 x 8
  `Molecular Type`  `X-ray`    EM    NMR `Multiple methods` Neutron Other  Total
  <chr>            <dbl> <dbl> <dbl>      <dbl>    <dbl> <dbl> <dbl>
1 Protein (only)    169563 16774 12578      208      81    32 199236
2 Protein/Oligosacc~ 9939 2839 34      8        2    0 12822
3 Protein/NA        8801 5062 286      7        0    0 14156
4 Nucleic acid (onl~ 2890 151 1521     14        3    1 4580
5 Other             170 10 33      0        0    0 213
6 Oligosaccharide (~ 11 0 6      1        0    4 22
```

I want to clean the column name so they are all lower case and don't have spaces in them

```
colnames(pdbstats)
```

```
[1] "Molecular Type"  "X-ray"          "EM"             "NMR"
[5] "Multiple methods" "Neutron"        "Other"          "Total"
```

```
library(janitor)
```

Attaching package: 'janitor'

The following objects are masked from 'package:stats':

```
chisq.test, fisher.test
```

```
df <- clean_names(pdbstats)
df
```

```
# A tibble: 6 x 8
  molecular_type      x_ray    em    nmr multiple_methods neutron other  total
  <chr>            <dbl> <dbl> <dbl>      <dbl>    <dbl> <dbl> <dbl>
1 Protein (only)    169563 16774 12578      208      81    32 199236
2 Protein/Oligosacchar~ 9939 2839 34      8        2    0 12822
3 Protein/NA        8801 5062 286      7        0    0 14156
4 Nucleic acid (only)  2890 151 1521     14        3    1 4580
5 Other             170 10 33      0        0    0 213
6 Oligosaccharide (onl~ 11 0 6      1        0    4 22
```

Total number of X-ray structures

```
sum(df$x_ray)
```

```
[1] 191374
```

Percentage of structures in the PDB are solved by X-ray

```
sum(df$x_ray)/sum(df$total) * 100
```

```
[1] 82.83549
```

Total number of EM (Electron Microscopy)

```
sum(df$em)/sum(df$total) * 100
```

```
[1] 10.75017
```

Total number of structures

```
sum(df$total)
```

```
[1] 231029
```

Q2: What proportion of structures in the PDB are protein? A2: 0.8623852

```
( df[1, "total"])/sum(df$total)
```

```
total  
1 0.8623852
```

The main Mol* homepage at: <https://molstar.org/viewer/>. We can input our own PDB files or just give it a PDB database accession code (4 letter PDB code)




Figure 1: Molecular view of 1HSG



Figure 2: Clear Ligand




Figure 3: Water 308 Bonding




Figure 4: No Polymer




Figure 5: A&B Residues of Aspartate(ASP25 Amino Acids)

Q4: Water molecules normally have 3 atoms. Why do we see just one atom per water molecule in this structure? A4: In this case, the hydrogen atoms are attached at a certain angle that makes them not visible and makes the molecule polar.

Q5: There is a critical “conserved” water molecule in the binding site. Can you identify this water molecule? What residue number does this water molecule have

A5: Yes, it is water 308.

Q6: Generate and save a figure clearly showing the two distinct chains of HIV-protease along with the ligand. You might also consider showing the catalytic residues ASP 25 in each chain and the critical water (we recommend “Ball & Stick” for these side-chains). Add this figure to your Quarto document.

3. Intro to Bio3D in R

We can use the **bio3D** package for structural bioinformatics to read PDB data into R

```
library(bio3d)

pdb <- read.pdb("1hsg")
```

Note: Accessing on-line PDB file

```
pdb
```

```
Call: read.pdb(file = "1hsg")
```

```
Total Models#: 1
Total Atoms#: 1686, XYZs#: 5058 Chains#: 2 (values: A B)

Protein Atoms#: 1514 (residues/Calpha atoms#: 198)
Nucleic acid Atoms#: 0 (residues/phosphate atoms#: 0)

Non-protein/nucleic Atoms#: 172 (residues: 128)
Non-protein/nucleic resid values: [ HOH (127), MK1 (1) ]
```

Protein sequence:

```
PQITLWQRPLVTIKIGGQLKEALLDTGADDTVLEEMSLPGRWKPKMIGGIGGFIKVRQYD
QILIEICGHKAIGTVLVGPTPVNIIGRNLLTQIGCTLNFPQITLWQRPLVTIKIGGQLKE
ALLDTGADDTVLEEMSLPGRWKPKMIGGIGGFIKVRQYDQILIEICGHKAIGTVLVGPTP
VNIIGRNLLTQIGCTLNF
```

```
+ attr: atom, xyz, seqres, helix, sheet,
      calpha, remark, call
```

Q7: How many amino acid residues are there in this pdb object? A7: There are 198 amino acid residues.

```
length(pdbseq(pdb))
```

```
[1] 198
```

Q8: Name one of the two non-protein residues? A8: MK1

Q9: How many protein chains are in this structure? A9: There are 2, chains A and B

Looking at the 'pdb()' object in more detail

```
attributes(pdb)
```

```
$names
```

```
[1] "atom" "xyz" "seqres" "helix" "sheet" "calpha" "remark" "call"
```

```
$class
```

```
[1] "pdb" "sse"
```

```
head(pdb$atom)
```

	type	eleno	elety	alt	resid	chain	resno	insert	x	y	z	o	b
1	ATOM	1	N	<NA>	PRO	A	1	<NA>	29.361	39.686	5.862	1	38.10
2	ATOM	2	CA	<NA>	PRO	A	1	<NA>	30.307	38.663	5.319	1	40.62
3	ATOM	3	C	<NA>	PRO	A	1	<NA>	29.760	38.071	4.022	1	42.64
4	ATOM	4	O	<NA>	PRO	A	1	<NA>	28.600	38.302	3.676	1	43.40
5	ATOM	5	CB	<NA>	PRO	A	1	<NA>	30.508	37.541	6.342	1	37.87
6	ATOM	6	CG	<NA>	PRO	A	1	<NA>	29.296	37.591	7.162	1	38.40

	segid	elesy	charge
1	<NA>	N	<NA>
2	<NA>	C	<NA>
3	<NA>	C	<NA>
4	<NA>	O	<NA>
5	<NA>	C	<NA>
6	<NA>	C	<NA>

Let's try a new function not yet in the bio3d package. It requires the **r3dmol** package that we need to install with 'install.packages("r3dmol")' and 'install.packages("shiny")'

```
#source("https://tinyurl.com/viewpdb")  
#view.pdb(pdb, backgroundColor = "white")
```

Q10. Which of the packages above is found only on BioConductor and not CRAN?

A10: BiocManager

Q11. Which of the above packages is not found on BioConductor or CRAN?: A11. The bio3d package.

Q12. True or False? Functions from the devtools package can be used to install packages from GitHub and BitBucket? A12. True

4. Predicting functional dynamics

We can use the 'nma()' function in bio3d to predict the large-scale functional motion of biomolecules.

```
adk <- read.pdb("6s36")
```

Note: Accessing on-line PDB file
PDB has ALT records, taking A only, rm.alt=TRUE

```
adk
```

Call: read.pdb(file = "6s36")

```
Total Models#: 1
  Total Atoms#: 1898, XYZs#: 5694 Chains#: 1 (values: A)

Protein Atoms#: 1654 (residues/Calpha atoms#: 214)
Nucleic acid Atoms#: 0 (residues/phosphate atoms#: 0)

Non-protein/nucleic Atoms#: 244 (residues: 244)
Non-protein/nucleic resid values: [ CL (3), HOH (238), MG (2), NA (1) ]
```

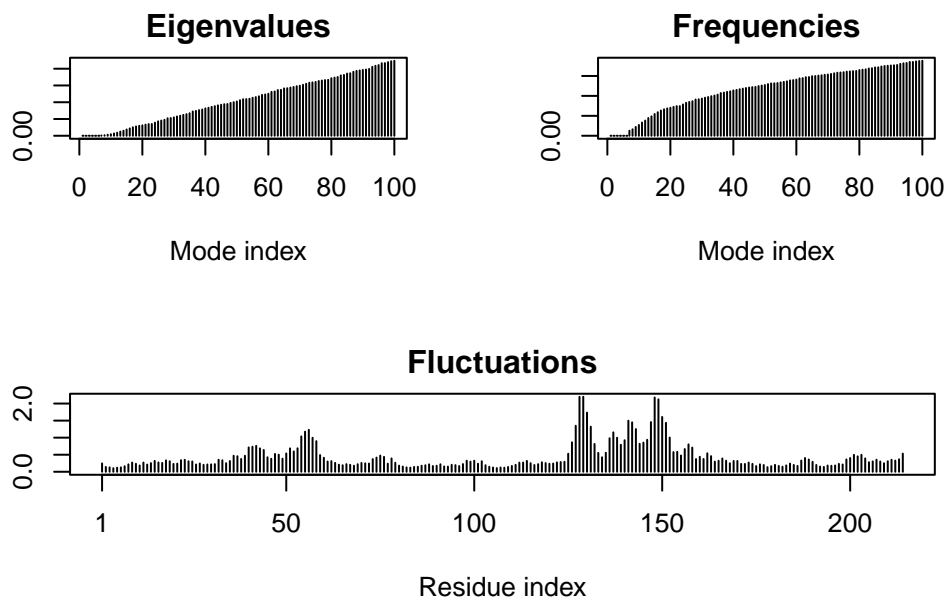
```
Protein sequence:
MRIILLGAPGAGKGTQAQFIMEKYGIPQISTGDMRLRAAVKSGSELGKQAKDIMDAGKLV
DELVIALVKERIAQEDCRNGFLLDGFPRTPQADAMKEAGINVDYVLEFDVPDELIVDKI
VGRRVHAPSGRVYHVKFNPVKVEGKDDVTGEELTTRKDDQEETVRKRLVEYHQM TAPLIG
YYSKEAEAGNTKYAKVDGTPVAEVRADLEKILG
```

```
+ attr: atom, xyz, seqres, helix, sheet,
      calpha, remark, call
```

```
m <- nma(adk)
```

```
Building Hessian...      Done in 0.01 seconds.
Diagonalizing Hessian... Done in 0.22 seconds.
```

```
plot(m)
```



Write out a trajectory of the predicted molecular motion:

```
meow <- mktrj(m, file="adk_m7.pdb")
```

Q13. How many amino acids are in this sequence, i.e. how long is this sequence?

A13. This sequence is 214 amino acids long.

```
library(bio3d)
aa <- get.seq("1ake_A")
```

Warning in get.seq("1ake_A"): Removing existing file: seqs.fasta

Fetching... Please wait. Done.

```
aa
```

```

      1      .      .      .      .      .      60
pdb|1AKE|A  MRIILLGAPGAGKGTQAQFIMEKYGIPQISTGDMRLRAAVKSGSELGKQAKDIMDAGKLV
      1      .      .      .      .      .      60
      61      .      .      .      .      .      120
pdb|1AKE|A  DELVIALVKERIAQEDCRNGFLLDGFPRPTIPQADAMKEAGINVDYVLEFDVPDELIVDRI

```

```

        61      .      .      .      .      .      .      120
        121     .      .      .      .      .      .      180
pdb|1AKE|A  VGRRVHAPSGRVYHVKNPPKVEGKDDVTGEELTTRKDDQEETVRKRLVEYHQMTAPLIG
        121     .      .      .      .      .      .      180

        181     .      .      .      214
pdb|1AKE|A  YYSKEAEAGNTKYAKVDGTPVAEVRADLEKILG
        181     .      .      .      214

```

Call:

```
read.fasta(file = outfile)
```

Class:

```
fasta
```

Alignment dimensions:

```
1 sequence rows; 214 position columns (214 non-gap, 0 gap)
```

```
+ attr: id, ali, call
```